

## Problem 2

### Convolutional Layer Output size calculation

Received [32, 16, 16, 16] and expected [32, 16, 16, 16]

#### 1. Convolutional Layer

##### a. Forward :

Received output shape: (1, 4, 4, 2), Expected output shape: (1, 4, 4, 2)  
Difference: 5.110565335399418e-08

##### b. Backward

dimg Error: 1.6670134158491315e-08  
dw Error: 1.056723541487401e-08  
db Error: 1.1157795914175144e-10  
dimg Shape: (15, 8, 8, 3) (15, 8, 8, 3)

#### 2. Pooling Layer (Max)

##### a. Forward

Received output shape: (1, 3, 3, 1), Expected output shape: (1, 3, 3, 1)  
Difference: 1.8750000280978013e-08

##### b. Backward

dimg Error: 3.276186843072994e-12  
dimg Shape: (15, 8, 8, 3) (15, 8, 8, 3)

#### 3. Test Small Convolutional Neural Network

Testing initialization ...

Passed!

Testing test-time forward pass ...

Passed!

Testing the loss ...

Passed!

Testing the gradients (error should be no larger than 1e-6) ...

conv1\_b relative error: 1.01e-09

conv1\_w relative error: 1.05e-09

fc1\_b relative error: 3.65e-10

fc1\_w relative error: 3.95e-07

Param names : Conv1\_w, Conv1\_b, fc1\_w, fc1\_b

#### 4. Train a Network : Conv->Max Pool->flatten->FC->GeLU->FC

##### Input Data :

Data shape: (40000, 32, 32, 3)  
Flattened data input size: 3072  
Number of data classes: 20

##### Layers defined:

```
ConvLayer2D(input_channels=3, kernel_size=5, number_filters=36, stride=2,
padding=1, name="conv1", init_scale=.02)
MaxPoolingLayer(pool_size=2, stride=2, name="maxpool1"),
flatten(name="flatten"),
fc(1764, 48, 2e-2, name="fc1"),
gelu(name="gelu1"),
fc(48, 20, 2e-2, name="fc2")
```

(Iteration 1 / 20000) Average loss: 3.001034925522668

100%|██████████| 4000/4000 [2:04:13<00:00, 1.86s/it]

(Epoch 1 / 5) Training Accuracy: 0.36765, Validation Accuracy: 0.3423

0%| | 1/4000 [00:01<2:00:46, 1.81s/it]

(Iteration 4001 / 20000) Average loss: 2.3154061446075054

40%|██████ | 1584/4000 [1:11:03<1:08:09, 1.69s/it]

(Epoch 2 / 5) Training Accuracy: 0.422725, Validation Accuracy: 0.3765

0%| | 1/4000 [00:01<1:56:45, 1.75s/it]

(Iteration 8001 / 20000) Average loss: 1.9999308835295233

100%|██████████| 4000/4000 [1:50:08<00:00, 1.65s/it]

(Epoch 3 / 5) Training Accuracy: 0.471775, Validation Accuracy: 0.4034

0%| | 1/4000 [00:01<1:44:44, 1.57s/it]

(Iteration 12001 / 20000) Average loss: 1.867676979601489

100%|██████████| 4000/4000 [1:44:30<00:00, 1.57s/it]


(Epoch 4 / 5) Training Accuracy: 0.47985, Validation Accuracy: 0.391

0%| | 1/4000 [00:01<1:47:02, 1.61s/it]


(Iteration 16001 / 20000) Average loss: 1.769121363689662

100%|██████████| 4000/4000 [1:49:07<00:00, 1.64s/it]

(Epoch 5 / 5) Training Accuracy: 0.522925, Validation Accuracy: 0.415



## 5. Visualization Layer



## 6. Visualization Inline Question :

The first layer of the model is defined as follows :

```
- (C) input_channels = 3
```

- (K) kernel\_size = 5
- (N) number\_filters=36
- (S) stride=2
- (P) padding=1
- Input image size = (32,32,3) => (input\_height, input\_width, channels)
- filter size = (5,5,3,36) => (K,K,C,N)

The above visualization depicts :

- 6x6 grid representing 36 filter
- Each small grid i.e. 5 x 5 map represents a color (r,g,b) matrix applied to the input image.
- One square represents

Analyses of the above visualization :

- We can see that in some cases, the filter is the similar across the channels (the first row), and in others, the filters differ (the last row)
- The darker squares indicate smaller weights and the lighter squares represent larger weights.
- We can notice the top left denotes the color filter RED (r,g,b) = (0-255,0,0), top right GREEN (r,g,b) = (0,0-255,0) and similarly BLUE.
- Hence, the filters on the first row detect a gradient from redder shades in the top left to dark blacker in the bottom right.

## 7. Extra Credit

### ### EXTRA CREDIT Inline Answer:

- Plotted the [confusion matrix] ([https://en.wikipedia.org/wiki/Confusion\\_matrix](https://en.wikipedia.org/wiki/Confusion_matrix)) of model's predictions on the test set.
- Justified different trends to see which classes are frequently misclassified as other classes (e.g. are the two vehicle superclasses frequently confused with each other?)
  - We can see in the below graph x-axis denotes expected classes and y denotes predicted classes.
  - Square which are darker are the ones that hold more classification weight for that particular (row-col) predicted-expected mapping. In ideal case along the diagonal the values would be the darkest meaning correct classifications.
  - We mark the trend where there are darker blocks in the heat map of confusion matrix, accounting for misclassifications. Examples of misclassification trends are as follows :
    - 2, 18 i.e trees and flower (Because of the possible color matrix being similar)
    - Similarly 18,19 classes i.e. vehicle\_1 and vehicle\_2 are often mixed and misclassified by the model, possibly because of object similarity
    - Similarly large\_natural\_outdoor\_scenes and aquatic mammals are misclassified, possibly because of close color patterns and arrangements (Attached example after confusion matrix)
- Also plotted a pair plot between samples to check for any correlation between them.

